

Loss Function Quiz – Questions and Answers

Q: What is the primary purpose of a loss function in training ML models?

A: A loss function measures how far the model's predictions deviate from the true target values.

Q: How does a loss function quantitatively measure the performance of an ML model?

A: It provides a numerical value that quantifies model prediction errors.

Q: What are the two main types of target variables mentioned in the context of loss functions?

A: The two main target variable types are continuous (regression) and categorical (classification).

Q: What must be done to the loss function to obtain the best/optimal set of model parameters?

A: To get optimal model parameters, the loss function must be minimized.

Q: What are the two alternative names for the Mean Squared Error (MSE)?

A: Mean Squared Error (MSE) is also called L2 Loss or Quadratic Loss.

Q: What is the formula for L_MSE?

A: $L_{MSE} = (1/n) \sum (y_i - \hat{y}_i)^2$.

Q: Why is MSE a very widely used loss function?

A: MSE is widely used because it is smooth, differentiable, and emphasizes larger errors.

Q: How does MSE's penalty mechanism affect larger errors?

A: MSE penalizes larger errors quadratically, increasing their influence.

Q: Due to its penalty mechanism, what is MSE sensitive to in a dataset?

A: MSE is sensitive to outliers because squared terms amplify extreme deviations.

Q: What property of the MSE function guarantees a global minimum for linear models?

A: Its convex nature ensures the existence of a single global minimum.

Q: What is the definition and unit relationship of Root Mean Squared Error (RMSE)?

A: RMSE is the square root of MSE and shares the same units as the target variable.

Q: Statistically, minimizing MSE is equivalent to obtaining which estimation parameters?

A: Minimizing MSE equals finding maximum likelihood estimates under Gaussian noise.

Q: Write out the normal distribution probability density function for noise ϵ_i .

A: $p(y_i | \hat{y}_i) = 1/\sqrt{2\pi\sigma^2} * \exp(-(y_i - \hat{y}_i)^2 / 2\sigma^2)$.

Q: Based on the log-likelihood function derived, what specific term is minimized?

A: The sum of squared errors term corresponds to the MSE loss being minimized.

Q: Under the Gaussian noise assumption, why does it make sense to use MSE?

A: MSE is justified when data errors follow a normal (Gaussian) distribution.

Q: What is the alternative name for the Mean Absolute Error (MAE)?

A: Mean Absolute Error (MAE) is also known as L1 Loss or Mean Absolute Deviation.

Q: What is the formula for L_MAE?

A: $L_{MAE} = (1/n) \sum |y_i - \hat{y}_i|$.

Q: How does the robustness of MAE compare to MSE?

A: MAE is more robust to outliers than MSE because it applies linear penalties.

Q: What specific characteristic of $|x|$ is illustrated in the graph?

A: The $|x|$ graph forms a V-shape with a sharp vertex at zero.

Q: What is the practical consequence of the derivative discontinuity for MAE?

A: The discontinuous derivative at zero causes unstable gradients in optimization.

Q: Explain the core difference between the L1 Loss and L2 Loss.

A: L1 uses absolute errors while L2 squares them, making L2 more sensitive to large deviations.

Q: Which loss function would generally be preferred if a dataset contains many outliers?

A: MAE is preferred when datasets contain extreme outliers due to its linear error handling.

Q: Why is easy differentiability important for a loss function?

A: Differentiability allows gradient-based algorithms to optimize parameters effectively.

Q: If prediction errors are Gaussian, which loss function is justified?

A: MSE suits models assuming Gaussian-distributed errors.

Q: What is the relationship between maximizing log-likelihood and minimizing loss?

A: Maximizing log-likelihood equals minimizing its negative, the loss function.

Q: What mathematical property of the logarithm simplifies optimization?

A: The log converts products into sums, simplifying optimization.

Q: What is the significance of the $(n/2)\log(2\pi\sigma^2)$ term?

A: The constant term $(n/2)\log(2\pi\sigma^2)$ does not affect minimization.

Q: What does ϵ_i represent in $y_i = \mu_i + \epsilon_i$?

A: ϵ_i represents random noise or deviation in observed data.

Q: What is the mathematical definition of the likelihood function L?

A: $L = \prod p(y_i | \mu_i)$ defines the likelihood under i.i.d. assumption.

Q: When is it guaranteed that a local minimum is also global?

A: For convex losses like MSE, any local minimum is also the global minimum.